

Project By: Trapti Khandelwal | Mohnish Lavania | Sri Lalana | Rahul Singh



REAL TIME OBJECT DETECTION (BASIC-I)

GUIDED BY: VEDANT DWIVEDI | SOMESH CHOUDHARY

Report

Introduction

Making a general object detection model(pretrained) such that it can be used in flutter applications and other web application calling the model via api(restful).

Object detection models seek to identify the presence of relevant objects in images and classify those objects into relevant classes.

- Basic Approach:- Our basic approach was converting a pre-trained yolov5 object detection model .

What is YOLO Object Detection?

YOLO (“You Only Look Once”) is an effective real-time object recognition algorithm, first described in the seminal 2015 paper by Joseph Redmon et al. In this article we introduce the concept of object detection, the YOLO algorithm itself, and one of the algorithm’s open source implementations: Darknet.

Image classification is one of the many exciting applications of convolutional neural networks. Aside from simple image classification, there are plenty of fascinating problems in computer vision, with object detection being one of the most interesting. It is commonly associated with self-driving cars where systems blend computer vision, LIDAR and other technologies to generate a multidimensional representation of the road with all its participants. Object detection is also commonly used in video surveillance, especially in crowd monitoring to prevent terrorist attacks, count people for general statistics or analyze customer experience with walking paths within shopping centers.

Object Detection Overview

To explore the concept of object detection it is useful to begin with image classification. It goes through levels of incremental complexity.

Image classification (1) aims at assigning an image to one of a number of different categories (e.g. car, dog, cat, human, etc.), essentially answering the question “What is in this picture?”. One image has only one category assigned to it.

Object localization (2) then allows us to locate our object in the image, so our question changes to “What is it and where it is?”.

In a real real-life scenario, we need to go beyond locating just one object but rather multiple objects in one image. For example, a self-driving car has to find the location of other cars, traffic lights, signs, humans and to take appropriate action based on this information.

YOLO algorithm

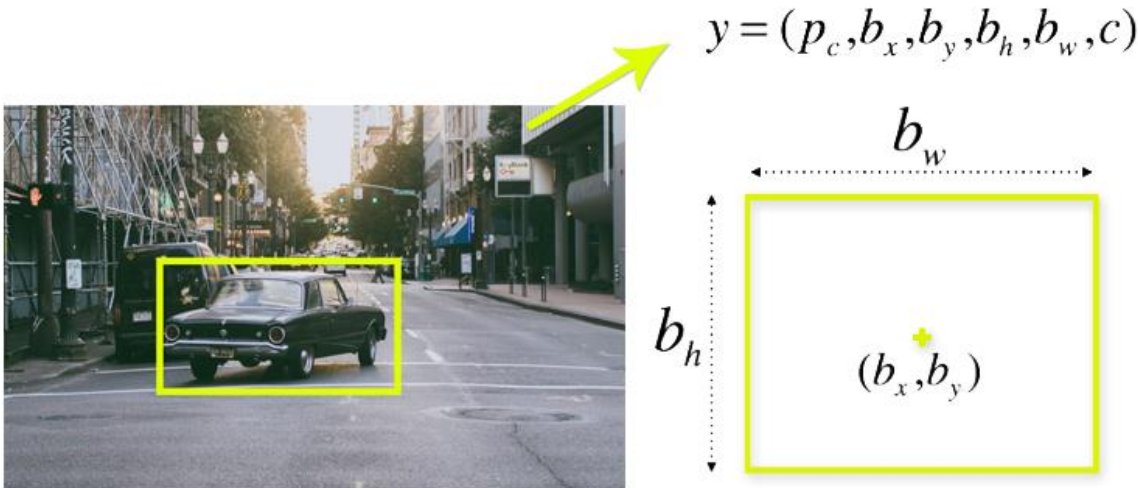
There are a few different algorithms for object detection and they can be split into two groups:

- Algorithms based on classification. They are implemented in two stages. First, they select regions of interest in an image. Second, they classify these regions using convolutional neural networks. This solution can be slow because we have to run predictions for every selected region. A widely known example of this type of algorithm is the Region-based convolutional neural network (RCNN) and its cousins Fast-RCNN, Faster-RCNN and the latest addition to the family: Mask-RCNN. Another example is RetinaNet.
- Algorithms based on regression – instead of selecting interesting parts of an image, they predict classes and bounding boxes for the whole image in one run of the algorithm. The two best known examples from this group are the YOLO (You Only Look Once) family algorithms and SSD (Single Shot Multibox Detector). They are commonly used for real-time object detection as, in general, they trade a bit of accuracy for large improvements in speed.

To understand the YOLO algorithm, it is necessary to establish what is actually being predicted. Ultimately, we aim to predict a class of an object and the bounding box specifying object location. Each bounding box can be described using four descriptors:

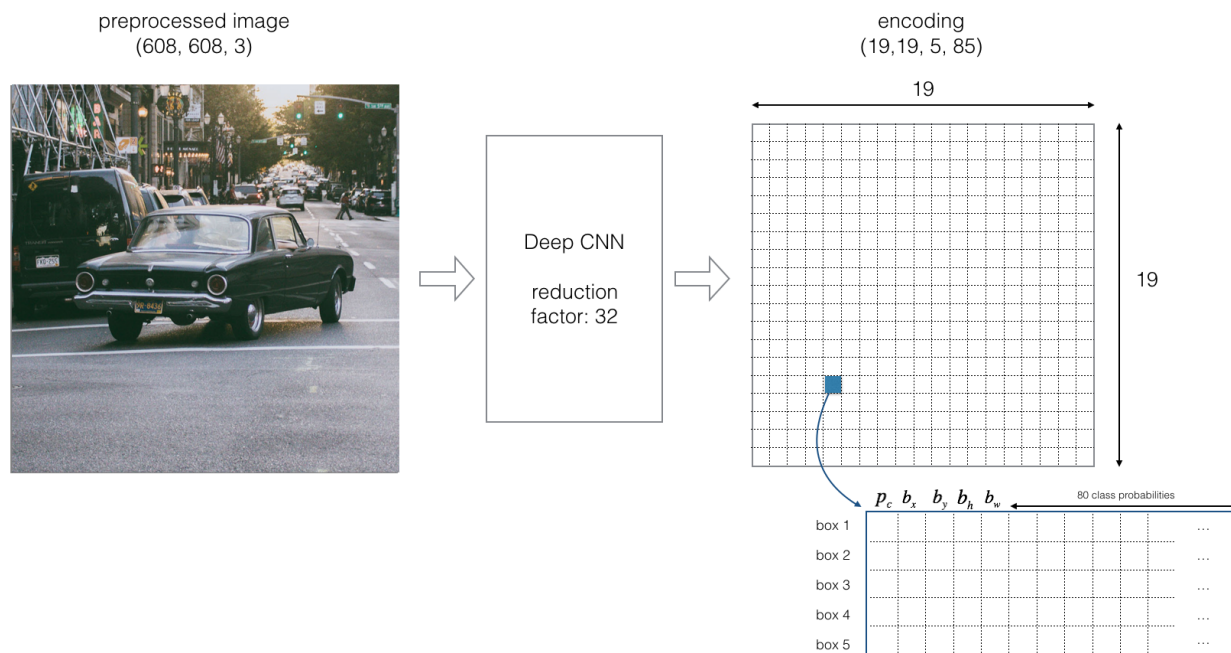
- center of a bounding box (bxby)
- width (bw)
- height (bh)
- value cis corresponding to a class of an object (such as: car, traffic lights, etc.).

In addition, we have to predict the pc value, which is the probability that there is an object in the boundingbox.



As we mentioned above, when working with the YOLO algorithm we are not searching for interesting regions in our image that could potentially contain an object.

Instead, we are splitting our image into cells, typically using a 19×19 grid. Each cell is responsible for predicting 5 bounding boxes (in case there is more than one object in this cell). Therefore, we arrive at a large number of 1805 bounding boxes for one image.



Most of these cells and bounding boxes will not contain an object. Therefore, we predict the value p_c , which serves to remove boxes with low object probability and bounding boxes with the highest shared area in a process called non-max suppression.

Darknet – a YOLO implementation

There are a few different implementations of the YOLO algorithm on the web. Darknet is one such open source neural network framework (a PyTorch implementation can be found [here](#) or with some extra fast.ai functionality [here](#); a Keras implementation can be found [here](#)). Darknet was written in the C Language and CUDA technology, which makes it really fast and provides for making computations on a GPU, which is essential for real-time predictions.

Installation is simple and requires running just 3 lines of code (in order to use GPU it is necessary to modify the settings in the Makefile script after cloning the repository).

After installation, we can use a pre-trained model or build a new one from scratch. For example here's how you can detect objects on your image using model pre-trained on COCO dataset:

Introduction to YOLO v5

YOLOv5 is a recent release of the YOLO family of models. YOLO v5 is the first of YOLO models to be written in the PyTorch framework and it is much more lightweight and easy to use. It is much easier to get started with and offers you much greater development speed when moving into deployment.

Installing the YOLOv5 Environment

To start off with YOLOv5 we first clone the YOLOv5 repository and install dependencies. This will set up our programming environment to be ready to running object detection training and inference commands.

```
!git clone https://github.com/ultralytics/yolov5 # clone repo
!pip install -U -r yolov5/requirements.txt # install dependencies%cd /content/yolov5
```

Then, we can take a look at our training environment provided to us for free from Google Colab.

```
import torch
from IPython.display import Image # for displaying images
from utils.google_utils import gdrive_download # for downloading models/datasetsprint('torch %s %s' % (torch.__version__,
torch.cuda.get_device_properties(0) if torch.cuda.is_available() else 'CPU'))
```

It is likely that you will receive a Tesla P100 GPU from Google Colab. Here is what We received:

```
torch 1.5.0+cu101 _CudaDeviceProperties(name='Tesla P100-PCIE-16GB', major=6, minor=0, total_memory=16280MB,
multi_processor_count=56)
```

The GPU will allow us to accelerate training time. Colab is also nice in that it come preinstalled with `torch` and `cuda`.

Download Custom YOLOv5 Object Detection Data

In this tutorial we will download custom object detection data in YOLOv5 format from Roboflow. You can follow along with the public blood cell dataset or upload your own dataset.

Once you have labeled data, to get move your data into Roboflow, create a free account and then you can drag your dataset in in any format: (VOC XML, COCO JSON, TensorFlow Object Detection CSV, etc).

Once uploaded you can choose preprocessing and augmentation steps:

The settings chosen for the BCCD example dataset

Then, click **Generate** and **Download** and you will be able to choose YOLOv5 PyTorch format.

Select “YOLO v5 PyTorch”

When prompted, be sure to select “Show Code Snippet.” This will output a download curl script so you can easily port your data into Colab in the proper format.

```
curl -L "https://public.roboflow.ai/ds/YOUR-LINK-HERE" > roboflow.zip; unzip roboflow.zip; rm roboflow.zip
```

Downloading in Colab...

Download a custom object detection dataset in YOLOv5 format

The export creates a YOLOv5 .yaml file called **data.yaml** specifying the location of a YOLOv5 **images** folder, a YOLOv5 **labels** folder, and information on our custom classes.

Define YOLOv5 Model Configuration and Architecture

Next we write a model configuration file for our custom object detector. For this tutorial, we chose the smallest, fastest base model of YOLOv5. You have the option to pick from other YOLOv5 models including:

- YOLOv5s
- YOLOv5m
- YOLOv5l
- YOLOv5x

You can also edit the structure of the network in this step, though rarely will you need to do this. Here is the YOLOv5 model configuration file, which we term `custom_yolov5s.yaml`:

```
nc: 3
depth_multiple: 0.33
width_multiple: 0.50anchors:
  - [10,13, 16,30, 33,23]
  - [30,61, 62,45, 59,119]
  - [116,90, 156,198, 373,326]backbone:
  [[-1, 1, Focus, [64, 3]],
  [-1, 1, Conv, [128, 3, 2]],
  [-1, 3, Bottleneck, [128]],
  [-1, 1, Conv, [256, 3, 2]],
  [-1, 9, BottleneckCSP, [256]],
  [-1, 1, Conv, [512, 3, 2]],
  [-1, 9, BottleneckCSP, [512]],
  [-1, 1, Conv, [1024, 3, 2]],
  [-1, 1, SPP, [1024, [5, 9, 13]]],
  [-1, 6, BottleneckCSP, [1024]],
  ]head:
  [[-1, 3, BottleneckCSP, [1024, False]],
  [-1, 1, nn.Conv2d, [na * (nc + 5), 1, 1, 0]],
  [-2, 1, nn.Upsample, [None, 2, "nearest"]],
  [[-1, 6], 1, Concat, [1]],
  [-1, 1, Conv, [512, 1, 1]],
  [-1, 3, BottleneckCSP, [512, False]],
  [-1, 1, nn.Conv2d, [na * (nc + 5), 1, 1, 0]],
  [-2, 1, nn.Upsample, [None, 2, "nearest"]],
  [[-1, 4], 1, Concat, [1]],
  [-1, 1, Conv, [256, 1, 1]],
  [-1, 3, BottleneckCSP, [256, False]],
  [-1, 1, nn.Conv2d, [na * (nc + 5), 1, 1, 0]],[], 1, Detect, [nc, anchors]],
  ]
```

Training Custom YOLOv5 Detector

With our `data.yaml` and `custom_yolov5s.yaml` files ready to go we are ready to train!

To kick off training we running the training command with the following options:

- `img`: define input image size
- `batch`: determine batch size
- `epochs`: define the number of training epochs. (Note: often, 3000+ are common here!)
- `data`: set the path to our `yaml` file
- `cfg`: specify our model configuration
- `weights`: specify a custom path to weights. (Note: you can download weights from the Ultralytics Google Drive folder)
- `name`: result names
- `nosave`: only save the final checkpoint
- `cache`: cache images for faster training

And run the training command:

Training a custom YOLOv5 detector. It trains quickly!

During training, you want to be watching the `mAP@0.5` to see how your detector is learning to detect on your validation set, higher is better!

Evaluate Custom YOLOv5 Detector Performance

Now that we have completed training, we can evaluate how well the training procedure performed by looking at the validation metrics. The training script will drop tensorboard logs in `runs`. We visualize those here:

Visualizing tensorboard results on our custom dataset. And if you can't visualize Tensorboard for whatever reason the results can also be plotted with `utils.plot_results` and saving a `result.png`. Training plots in `.png` format

Run YOLOv5 Inference on Test Images

Now we take our trained model and make inference on test images. After training has completed model weights will save in `weights/`. For inference we invoke those weights along with a `conf` specifying model confidence (higher confidence required makes less predictions), and a inference `source`. `source` can accept a directory of images, individual images, video files, and also a device's webcam port. For source, I have moved our `test/*.jpg` to `test_infer/`.

```
!python detect.py --weights weights/last_yolov5s_custom.pt --img 416 --conf 0.4 --source ../test_infer
```

The inference time is extremely fast. On our Tesla P100, the YOLOv5s is hitting 7ms per image. This bodes well for deploying to a smaller GPU like a Jetson Nano (which costs only \$100).

Inference on YOLOv5s occurring at 142 FPS (.007s/image)

Finally, we visualize our detectors inferences on test images.

YOLOv5 inference on test images. It can also easily infer on video and webcam.

Recall and Precision

Average Precision	(AP) @[IoU=0.50:0.95	area= all	maxDets=100] = 0.352
Average Precision	(AP) @[IoU=0.50	area= all	maxDets=100] = 0.544
Average Precision	(AP) @[IoU=0.75	area= all	maxDets=100] = 0.378
Average Precision	(AP) @[IoU=0.50:0.95	area= small	maxDets=100] = 0.187
Average Precision	(AP) @[IoU=0.50:0.95	area=medium	maxDets=100] = 0.397
Average Precision	(AP) @[IoU=0.50:0.95	area= large	maxDets=100] = 0.459
Average Recall	(AR) @[IoU=0.50:0.95	area= all	maxDets= 1] = 0.296
Average Recall	(AR) @[IoU=0.50:0.95	area= all	maxDets= 10] = 0.496
Average Recall	(AR) @[IoU=0.50:0.95	area= all	maxDets=100] = 0.557
Average Recall	(AR) @[IoU=0.50:0.95	area= small	maxDets=100] = 0.358
Average Recall	(AR) @[IoU=0.50:0.95	area=medium	maxDets=100] = 0.619
Average Recall	(AR) @[IoU=0.50:0.95	area= large	maxDets=100] = 0.700

Average Precision	(AP) @[IoU=0.50:0.95	area= all	maxDets=100] = 0.470
Average Precision	(AP) @[IoU=0.50	area= all	maxDets=100] = 0.659
Average Precision	(AP) @[IoU=0.75	area= all	maxDets=100] = 0.515
Average Precision	(AP) @[IoU=0.50:0.95	area= small	maxDets=100] = 0.310
Average Precision	(AP) @[IoU=0.50:0.95	area=medium	maxDets=100] = 0.516
Average Precision	(AP) @[IoU=0.50:0.95	area= large	maxDets=100] = 0.610
Average Recall	(AR) @[IoU=0.50:0.95	area= all	maxDets= 1] = 0.362
Average Recall	(AR) @[IoU=0.50:0.95	area= all	maxDets= 10] = 0.597
Average Recall	(AR) @[IoU=0.50:0.95	area= all	maxDets=100] = 0.656
Average Recall	(AR) @[IoU=0.50:0.95	area= small	maxDets=100] = 0.504
Average Recall	(AR) @[IoU=0.50:0.95	area=medium	maxDets=100] = 0.704
Average Recall	(AR) @[IoU=0.50:0.95	area= large	maxDets=100] = 0.787

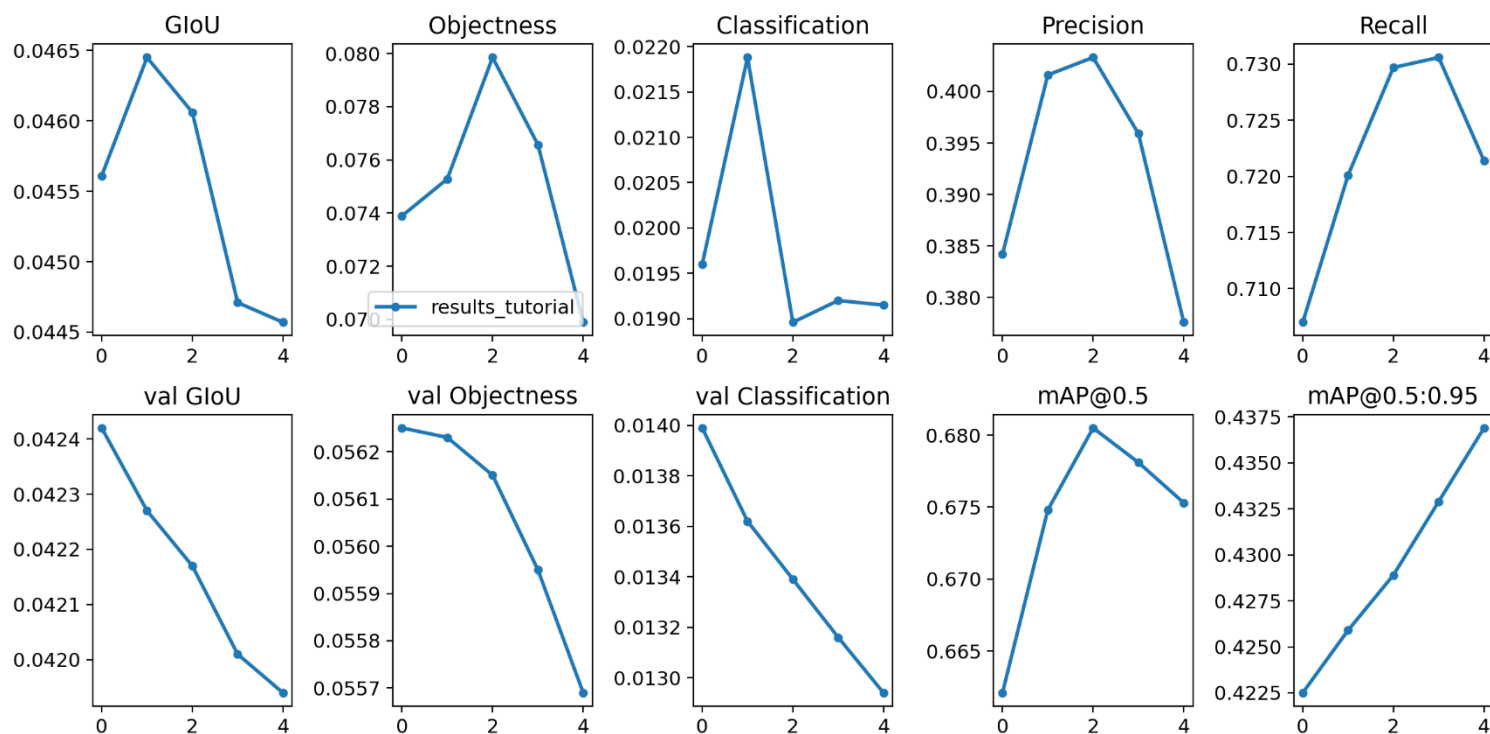
Model Summary:

```
from n  params module arguments
0      -1 1    3520 models.common.Focus [3, 32, 3]
1      -1 1   18560 models.common.Conv [32, 64, 3, 2]
2      -1 1   20672 models.common.Bottleneck [64, 64]
3      -1 1   73984 models.common.Conv [64, 128, 3, 2]
4      -1 1  161152 models.common.BottleneckCSP [128, 128, 3]
5      -1 1  295424 models.common.Conv [128, 256, 3, 2]
6      -1 1  641792 models.common.BottleneckCSP [256, 256, 3]
7      -1 1 1180672 models.common.Conv [256, 512, 3, 2]
8      -1 1  656896 models.common.SPP [512, 512, [5, 9, 13]]
9      -1 1 1905152 models.common.BottleneckCSP [512, 512, 2]
10     -1 1 1248768 models.common.BottleneckCSP [512, 512, 1, False]
11     -1 1  130815 torch.nn.modules.conv.Conv2d [512, 255, 1, 1]
12     -2 1      0 torch.nn.modules.upsampling.Upsample [None, 2, 'nearest']
13     [-1, 6] 1      0 models.common.Concat [1]
14     -1 1  197120 models.common.Conv [768, 256, 1, 1]
15     -1 1  313088 models.common.BottleneckCSP [256, 256, 1, False]
16     -1 1   65535 torch.nn.modules.conv.Conv2d [256, 255, 1, 1]
17     -2 1      0 torch.nn.modules.upsampling.Upsample [None, 2, 'nearest']
18     [-1, 4] 1      0 models.common.Concat [1]
19     -1 1   49408 models.common.Conv [384, 128, 1, 1]
20     -1 1   78720 models.common.BottleneckCSP [128, 128, 1, False]
21     -1 1   32895 torch.nn.modules.conv.Conv2d [128, 255, 1, 1]
22     [-1, 16, 11] 1      0 models.yolo.Detect [80, [[10, 13, 16, 30, 33, 23], [30, 61, 62, 45, 59, 119], [116, 90,
Model Summary: 165 layers, 7.07417e+06 parameters, 7.07417e+06 gradients
```

Epoch	gpu_mem	GIoU	obj	cls	total	targets	img_size
0/4	7.3G	0.04561	0.07389	0.0196	0.1391	217	640: 100% 8/8 [00:15<00:00, 1.90s/it]
Class		Images	Targets		P	R	mAP@.5: .95: 100% 8/8 [00:14<00:00, 1.80s/it]
all		128	929		0.384	0.707	0.662 0.422
Epoch	gpu_mem	GIoU	obj	cls	total	targets	img_size
1/4	11.1G	0.04645	0.07528	0.02188	0.1436	152	640: 100% 8/8 [00:06<00:00, 1.16it/s]
Class		Images	Targets		P	R	mAP@.5: .95: 100% 8/8 [00:04<00:00, 1.84it/s]
all		128	929		0.402	0.72	0.675 0.426
Epoch	gpu_mem	GIoU	obj	cls	total	targets	img_size
2/4	11.1G	0.04606	0.07985	0.01896	0.1449	222	640: 100% 8/8 [00:06<00:00, 1.16it/s]
Class		Images	Targets		P	R	mAP@.5: .95: 100% 8/8 [00:04<00:00, 1.94it/s]
all		128	929		0.403	0.73	0.681 0.429
Epoch	gpu_mem	GIoU	obj	cls	total	targets	img_size
3/4	11.1G	0.04471	0.07657	0.0192	0.1405	215	640: 100% 8/8 [00:06<00:00, 1.15it/s]
Class		Images	Targets		P	R	mAP@.5: .95: 100% 8/8 [00:04<00:00, 1.89it/s]
all		128	929		0.396	0.731	0.678 0.433
Epoch	gpu_mem	GIoU	obj	cls	total	targets	img_size
4/4	11.1G	0.04457	0.06989	0.01915	0.1336	153	640: 100% 8/8 [00:07<00:00, 1.14it/s]
Class		Images	Targets		P	R	mAP@.5: .95: 100% 8/8 [00:04<00:00, 1.89it/s]
all		128	929		0.378	0.721	0.675 0.437

Optimizer stripped from weights/last_tutorial.pt
5 epochs completed in 0.021 hours.

Parameters



Output



Conclusion

YOLO v5 is lightweight and extremely easy to use. YOLO v5 trains quickly, inferences quickly, and performs well.